

Liana Valdes

Ph.D. Graduate
Knight Foundation School of Computing and Information Sciences
Florida International University (FIU)
Miami, FL, 33199

+1 (786) 665-5542
Email: lvald108@fiu.edu
Web: <https://lvald.netlify.app/>
GitHub: github.com/lia54
LinkedIn: linkedin.com/in/liana-valdes/

RESEARCH & PROFESSIONAL INTERESTS

Storage, Distributed Systems, Caching Algorithms, ML for Systems, Systems for ML, OS, Team Leadership, Innovation, Motivation, Empathy, Technology, Networking, Creativity, and Relationship Building.

EDUCATION

Florida International University	21 August 2017 - 16 December 2023
<i>Doctor of Philosophy in Computer Science</i>	Miami, FL
<i>Advisor: Eminent Scholar Chaired Professor Raju Rangaswami.</i>	GPA: 3.83/4.0

Graduate Relevant Courses: Operating Systems (COP-5614), Analysis of Algorithms (COT-6405), Theory of Computation (COT-5310), Computer Communication and Networking Technologies (TCN-5030), Introduction to Algorithms (COT-5407), Secure Application Programming (CEN-5079), Data Base Management Systems (COP-5725), Introduction to Machine Learning (CAP-5610), Advanced Software Engineering (CEN-5010), Data Visualization (CAP-5738).

Florida International University	21 August 2017 - 17 December 2022
<i>Master of Science in Computer Science</i>	Miami, FL

Technological University of Havana “José Antonio Echeverría”	1 September 2009 - 20 July 2014
<i>Bachelor of Science in Telecommunications and Electronics Engineering</i>	Havana, Cuba
<i>Faculty of Telecommunications Engineering (FIT).</i>	GPA: 4.48/5

Undergraduate Relevant Courses: Linear Algebra and Analytical Geometry, Drafting, Political Economics of Capitalism, Philosophy and Society, Mathematics I, Mathematics II, Molecular Physics and Mechanics, Computerized and Automated Office Systems (Elective), Programming I, Introduction to Engineering Project, Chemistry, Electric Circuits I, Electric Circuits II, Electromagnetism and Optics, Analog Electronics I, Communication Statistics, Quantum Physics, Materials and Components, Probability and Statistics, Advanced Programming (Elective), Programming II, Differential Equations and Series, Sociopolitical Theory, Electrical Circuits III, Digital Electronics I, Digital Electronics II, Analog Electronics II, Analog Electronics III, Principles of Communications I, Principles of Communications II, Virtual Instrumentation (Elective), Transmission Lines, Research Methodology, Numerical Methods, Microprocessors I, Social Issues in Science and Technology, Digital Signal Processing, Field Theory, Antennas, Optical Communications, Telephone Communication, Accounting and Finance, Principal of Communications III, Network Management (Elective), Microprocessors II, Radio Electronics I, Radio Electronics II, Radio Propagation, Networks 1, Networks 2, Network Simulators (Elective), Transmission System, Microprocessor Systems, Business Administration, Satellite Communications (Elective), Principles of Television, Network 3, and IP Telephone (Elective).

RESEARCH PROJECTS

Extending storage systems to meet cloud needs Seagate Tech. & SyLab (FIU)	August 2021 - Present
--	------------------------------

- Designed TxFuse, a novel architecture that uses feature plugins to enhance storage systems.
- Designed the transactional coupling and reliable notification techniques to develop plugins.
- Evaluated feature plugin prototypes based on system performance and development complexity.

A distributed cache for cloud data centers | SyLab, ModLab, DAMRL (FIU) & HASLab (UMinho) May 2020 - Present

- Designed CaaS, a novel, distributed, and generalized caching for computing infrastructures in data centers.
- Improved hit rate and the service's and microservice's reading and writing performance with cache replacement.
- Designed a simulation framework to model a CaaS system with different components.

ML systems for caches in the cloud | SyLab, ModLab, BioRG (FIU) & VISA (ASU) August 2018 - February 2021

- Analyzed production workloads from different cloud service providers.
- Designed ML algorithms LeCaR and CACHEUS that outperform classic cache replacement algorithms.
- Reduced the cost of accessing SSD-based caches by reducing the number of writes we performed on the devices.

PUBLICATIONS

Project Silica: Towards Sustainable Cloud Archival Storage in Glass. October 2023

Liana Valdes, in Acknowledgement, The 29th ACM Symposium on Operating Systems Principles, SOSP'23.

Infusing Pub-Sub Storage with Transactions July 2022

Liana V. Rodriguez, John Bent, Tim Shaffer, and Raju Rangaswami, 14th ACM Workshop, HotStorage'22.

Unifying the Data Center Caching Layer - Feasible? Profitable? July 2021

Liana V. Rodriguez, Alexis Gonzalez, Pratik Poudel, Raju Rangaswami, and Jason Liu, 13th ACM Workshop, HotStorage'21.

Learning Cache Replacement with CACHEUS February 2021

Liana V. Rodriguez, Farzana Yusuf, Steven Lyons, Eysler Paz, Raju Rangaswami, Jason Liu, Ming Zhao, and Giri Narasimhan, 19th USENIX Conference, FAST'21.

Driving Cache Replacement with ML-Based LeCaR July 2018

Giuseppe Vietri, Liana V. Rodriguez, Wendy A. Martinez, Steven Lyons, Jason Liu, Raju Rangaswami, Ming Zhao, Giri Narasimhan. USENIX Workshop, HotStorage'18.

PRESENTATIONS

CORTX and FDMI | CORTX Meet the Architect Series November 2022

Liana V. Rodriguez and John Bent.

Infusing Pub-Sub Storage with Transactions | HotStorage'22 Presentation July 2022

Liana V. Rodriguez, John Bent, Tim Shaffer, and Raju Rangaswami.

Unifying the Data Center Caching Layer - Feasible? Profitable? | HotStorage'21 Presentation July 2021

Liana V. Rodriguez, Alexis Gonzalez, Pratik Paudel, Raju Rangaswami and Jason Liu.

Learning Cache Replacement with CACHEUS | Poster Annual FIU SCIS Research Day October 2019

Liana V. Rodriguez, Farzana Yusuf, Steven Lyons, Eysler Paz, Raju Rangaswami, Jason Liu, Ming Zhao, and Giri Narasimhan.

Driving Cache Replacement with ML-Based LeCaR | HotStorage'18 Poster session July 2018

Giuseppe Vietri, Liana V. Rodriguez, Wendy A. Martinez, Steven Lyons, Jason Liu, Raju Rangaswami, Ming Zhao, and Giri Narasimhan.

EXPERIENCE

Graduate Research Assistant/GAANN Fellowship | FIU, FL, US August 2017 - December 2023

- Analyzed production storage workloads and discovered workload patterns called workload primitives. As a result, I studied how these workloads function in the real world using a classification based on the composition of primitive workload types. I computed metrics to characterize the production storage workloads. For a one-week trace from MSR, Cambridge, 3.29% of the requests were write operations, and the unique and reuse footprints were 89.24% and 2.22%, respectively. The research dives into the intricacies of the data and examines the effects of several factors, providing valuable insights. This approach has improved cache performance in novel cache replacement algorithms capable of adapting to the data and tailored solutions to meet the needs of the production storage workload.
- Designed ML systems for cache replacement, LeCaR, and CACHEUS, which enhance storage caching performance. When evaluating a sample trace from MSR, Cambridge, LeCaR, and CACHEUS algorithms outperform state-of-the-art

caches, including LRU and ARC, with hit-rate values of 93.61% and 94.38%, respectively.

- Designed CaaS, a distributed caching service that improves hit rate, service, and microservice reading and writing performance. CaaS outperforms an SSD-based local cache commonly deployed in the data centers. The hit rate of CaaS is 96% for a one-day trace from Cloud VPS provider, outperforming the local cache hit rate of 4%. In these systems, ensuring data consistency, scalability, high availability, disaster recovery, and fault tolerance are addressed to guarantee optimal performance and reliability.

Research Intern | Seagate Technology, Remote, US

August 2021 - December 2021

- Seagate, a leading data storage company, developed CORTX, an open-source object storage solution. The solution facilitates storing and managing massive amounts of unstructured data, such as media files, logs, and backups. CORTX has an integrated software stack comprising several components. I tested the CORTX software stack to ensure seamless integration and deployment.
- Tested the CORTX's interface, File Data Manipulation Interface (FDMI), which uses plugins to implement storage features. I developed plugins such as a word count program that counts the number of times a word appears in each S3 PUT request to the storage. This plugin is equivalent to a program that polls a directory every 3 seconds and prints the 30 most popular words within each file. This research contributes to the design and development of more robust and efficient storage systems that can meet the ever-evolving needs of the modern world.
- Conducted a performance study on Motr, a distributed object storage system with advanced features and capabilities deployed at the Jülich Supercomputing Center (JSC) in Germany. The study aimed to analyze the system's performance and evaluate its efficiency. We reported the maximum achievable network and device bandwidth in a Sage platform x86 Motr cluster with sixteen clients and seven server nodes; each server had three tiers (tier-1, tier-2, and tier-3) that have NVMe, SSD, and HDD storage devices, respectively. The results showed 99% and 86% bandwidth utilization for read and write operations in tier-2 compared to the theoretical system limit.

Research Intern | Microsoft Research (MSR), Cambridgeshire, UK

January 2020 - March 2020

- Helped develop the software stack for a storage system using cutting-edge quartz glass technology as part of Microsoft's Project Silica team. This innovative technology is specifically designed to cater to the requirements of Azure Cloud archival workloads, enabling efficient and cost-effective storage. Unlike traditional storage systems, Project Silica's approach involves storing vast amounts of cold data, creating tiny 'voxels' into quartz glass.
- Designed ML systems using Deep Learning and Unsupervised Learning that can identify anomalies when tested using data from the Warner Bros. movie "Superman" written on glass. These systems included isolation forests and encoders/decoders, CNN, and GAN as anomaly detection systems. These ML models can help identify unusual patterns in the image data, particularly fine-grained errors within one sector (XY plane) and coarse-grained errors within one track (Z direction).

Telematics Specialist B | Cuban Telecommunications Company S.A.

September 2015 - June 2016

- Maintained uninterrupted and efficient communication services, crucial for monitoring the core infrastructure of PSTN, ATM, DSL, and public and private VoIP telephone systems. This involves closely observing the hardware, software, and network elements that comprise these systems' backbone. Any potential issues were detected and resolved promptly by continuously monitoring performance metrics such as speed, latency, uptime, and number of lines interrupted.
- Monitored the security of these systems to prevent unauthorized access and data breaches. With a focus on proactive monitoring and maintenance, these critical communication systems were highly reliable and available, ensuring seamless communication for businesses and individuals.
- Helped conduct a network assessment to ensure the continuous deployment of the Fiber to the Home (FTTH) networks in Old Havana. The assessment helped to identify the best action plan for deploying OTN devices.
- Helped assess and enhance PSTN network infrastructure to meet the community's needs.

Research Intern/Apprenticeship | Cuban Radio and Television Broadcasting Company

March 2012 - July 2014

- Analyzed the board layout to modify the UHF and VHF communication modules from PAL to NTSC television standards. This modification improved the frequency of the television signal, resulting in a much more precise and higher-quality picture.
- Helped to design RF matching networks that use passive components to achieve maximum power transfer to a 50-ohm or a real-world load, such as a television antenna. This involves analyzing the properties of the RF signal and

determining the best arrangement of passive components to match the load's impedance.

- Simulated results obtained with MATLAB, virtual instrumentation of E-field sensors, and spectrum/network analyzers. The simulations allowed us to visualize the RF signal's properties, including frequency, amplitude, and phase.

HONORS & AWARDS

USENIX Student Travel Award, FAST'18 & FAST'19 & FAST'23.	Sep. 2018, Sep. 2019, and Sept. 2023
GAANN Fellowship from the U.S. Department of Education (ED).	Jan. 2022 & Jan. 2023
CMD-IT/ACM Richard Tapia Celebration of Diversity in Computing Conference Scholarship.	Sep. 2022 & Sep. 2024
Grace Hooper Celebration of Women in Computing FIU Scholarship.	Oct. 2019 & Sep. 2022
USENIX SREcon24 Europe/Middle East/Africa Award	Nov. 2024

LEADERSHIP & ACTIVITIES

- Selected to plan and organize an event for Seagate's global CORTX Hackathon, focusing on integration challenges, possible solutions, and open-source software.
- Selected as a participant and anonymous reviewer nationwide for a peer-reviewed journal article for the ACM Transactions on Storage (TOS), the premier journal for publishing advancements in data storage research.
- Selected presentation slides for an annual Women in Computer Science (WiCS) meeting, focusing on sharing experiences and strategies for achieving academic success in computer science.
- Participated in Upsilon Pi Epsilon (UPE) activities, including workshops, presentations, tutorials, and leadership events to develop leadership skills, build industry connections, and achieve academic success.
- Selected as a volunteer for an eight-month teaching program for the Google CS First program, focusing on coding skills and learning computer science concepts using Scratch for kids under eight.
- Selected as team leader for four Systems Research Laboratory's research projects at FIU.

TEACHING ASSISTANT

Operating Systems	August 2018 & January 2019
Florida International University	Miami, FL
Transmission Systems II	2011 - 2012
Technological University of Havana "José Antonio Echeverría"	Havana, Cuba
Analog Electronics I	2010 - 2011
Technological University of Havana "José Antonio Echeverría"	Havana, Cuba

TECHNICAL SKILLS

Skills: Cache, Storage Systems, Machine Learning, Reinforcement Learning, Deep Learning, Supervised Learning, Unsupervised Learning, Artificial Intelligence, Large Language Model, Distributed Systems, Distributed caching, QoS algorithms.

Languages: Python, C#, C, C++, Java, Go, R, R++, TeX, Visual Basic, High-Level Assembly (HLA).

Tools/Frameworks: Git, GitLab, Docker, VirtualBox, VMware Workstation, VsCode, Tableau, MATLAB, LabVIEW, LaTeX, PyTorch, TensorFlow, Scikit-learn.

LEADERSHIP SKILLS

Team Leadership, Relationship Building, Emotional Intelligence, Forward-thinking, Problem-solving, Empowerment, Adaptability, Motivation, Innovation, Inspiration, Active Listening, Empathy, Presentation, Authenticity, Honesty, Creativity, Integrity, Decision-making, Accountability, Strategic Thinking, Communication, Coaching, and Management.

CERTIFICATIONS

CompTIA A+ 220-801 Certification Exam.

CompTIA A+ 220-802 Certification Exam.

LANGUAGE

Spanish: Fluent.

English: Full Professional Proficiency.

HOBBIES

Music: Traditional Cuban Music, Country Music, Pop Music, Blues, and Jazz.

Poetry: Cuban Poetry, Latino-American Poetry, Spanish Poetry, and American Poetry.

Chess: Tournaments, Titles and Rankings.

Fitness: Physical Fitness, Exercise, Training, and Walking.

Animation: Computer Animation, Disney Animation, and Japanese Animation.