

Liana Valdes Rodriguez

Ph.D. Graduate

Knight Foundation School of Computing and Information Sciences

Florida International University

Miami, FL, 33199

Mobile: +1-786-665-5542

Email: lvald108@fiu.edu

Web: <https://lvald.netlify.app/>

GitHub: github.com/lia54, LinkedIn: [linkedin.com/in/liana-valdes](https://www.linkedin.com/in/liana-valdes)

RESEARCH & PROFESSIONAL INTERESTS

Storage, Distributed Systems, Caching Algorithms, ML for Systems, Systems for ML, OS, Management, Team Leadership, Innovation, Motivation, Empathy, Technology, Networking, Creativity, and Relationship Building.

EDUCATION

Florida International University (FIU)

21 August 2017 - 16 December 2023

Doctor of Philosophy in Computer Science (CS)

Miami, FL

Advisor: Eminent Scholar Chaired Professor Raju Rangaswami.

GPA: 3.83/4.0

Graduate Relevant Courses: Operating Systems (OS, COP-5614), Analysis of Algorithms (COT-6405), Theory of Computation (COT-5310), Computer Communication and Networking Technologies (TCN-5030), Introduction to Algorithms (COT-5407), Secure Application Programming (CEN-5079), Data Base Management Systems (DBMS, COP-5725), Introduction to Machine Learning (ML, CAP-5610), Advanced Software Engineering (CEN-5010), Data Visualization (CAP-5738).

Florida International University (FIU)

21 August 2017 - 17 December 2022

Master of Science in Computer Science (CS)

Miami, FL

Advisor: Eminent Scholar Chaired Professor Raju Rangaswami.

Technological University of Havana “José Antonio Echeverría” (CUJAE)

1 September 2009 - 28 July 2014

Bachelor of Science in Electronics and Telecommunications Engineering

Havana, Cuba

Faculty of Telecommunications Engineering (FIT).

GPA: 4.48/5

Undergraduate Relevant Courses: Linear Algebra and Analytical Geometry, Drafting, Political Economics of Capitalism, Philosophy and Society, Mathematics I, Mathematics II, Molecular Physics and Mechanics, Computerized and Automated Office Systems (Elective), Programming I, Introduction to Engineering Project, Chemistry, Electric Circuits I, Electric Circuits II, Electromagnetism and Optics, Analog Electronics I, Communication Statistics, Quantum Physics, Materials and Components, Probability and Statistics, Advanced Programming (Elective), Programming II, Differential Equations and Series, Sociopolitical Theory, Electrical Circuits III, Digital Electronics I, Digital Electronics II, Analog Electronics II, Analog Electronics III, Principles of Communications I, Principles of Communications II, Virtual Instrumentation (Elective), Transmission Lines, Research Methodology, Numerical Methods, Microprocessors I, Social Issues in Science and Technology, Digital Signal Processing, Field Theory, Antennas, Optical Communications, Telephone Communication, Accounting and Finance, Principal of Communications III, Network Management (Elective), Microprocessors II, Radio Electronics I, Radio Electronics II, Radio Propagation, Networks 1, Networks 2, Network Simulators (Elective), Transmission System, Microprocessor Systems, Business Administration, Satellite Communications (Elective), Principles of Television, Network 3, and IP Telephone (Elective).

RESEARCH PROJECTS

Extensible Distributed Storage Systems | Seagate Technology & SyLab (FIU)

August 2021

- Developed TxFuse, a novel architecture that uses feature plugins to enhance distributed storage systems.
- Created a feature plugin taxonomy based on transactional coupling and reliable notification.
- Evaluated feature plugin prototypes based on performance and development complexity.

Distributed Caches in Data Centers | SyLab, ModLab, DAMRL (FIU) & HASLab (UMinho)

May 2020

- Developed CaaS, a novel, distributed caching service for cloud storage production workloads.
- Improved hit rate and performance of read and write operations compared to back-end storage.
- Developed a simulation framework to model the CaaS system's various components.

Caching Algorithms for Storage Caches | SyLab, BioRG (FIU) & VISA (ASU)

August 2017

- Analyzed production storage workloads from different cloud service providers.
- Developed LeCaR and CACHEUS, novel ML caching algorithms that achieve higher hit rates for production workloads.
- Reduced the cost of accessing SSD-based caches by reducing the total of writes sent to the devices.

PUBLICATIONS

Project Silica: Towards Sustainable Cloud Archival Storage in Glass

October 2023

Liana Valdes, in *Acknowledgement*,
The 29th ACM Symposium on Operating Systems Principles, SOSP'23.

Infusing Pub-Sub Storage with Transactions

July 2022

Liana V. Rodriguez, John Bent, Tim Shaffer, and Raju Rangaswami, 14th
ACM Workshop, HotStorage'22.

Unifying the Data Center Caching Layer - Feasible? Profitable?

July 2021

Liana V. Rodriguez, Alexis Gonzalez, Pratik Poudel, Raju Rangaswami, and Jason Liu, 13th
ACM Workshop, HotStorage'21.

Learning Cache Replacement with CACHEUS

February 2021

Liana V. Rodriguez, Farzana Yusuf, Steven Lyons, Eysler Paz, Raju Rangaswami, Jason Liu, Ming Zhao, and Giri
Narasimhan, 19th USENIX Conference, FAST'21.

To Cache or Not to Cache

Steven Lyons, Liana V. Rodriguez, and Raju Rangaswami.

Driving Cache Replacement with ML-Based LeCaR

July 2018

Giuseppe Vietri, Liana V. Rodriguez, Wendy A. Martinez, Steven Lyons, Jason Liu, Raju Rangaswami, Ming Zhao, and
Giri Narasimhan, USENIX Workshop, HotStorage'18.

PRESENTATIONS

CORTX and FDMI | CORTX Meet the Architect Series

November 2022

Liana V. Rodriguez and John Bent.

Infusing Pub-Sub Storage with Transactions | HotStorage'22 Presentation

July 2022

Liana V. Rodriguez, John Bent, Tim Shaffer, and Raju Rangaswami.

Unifying the Data Center Caching Layer - Feasible? Profitable? | HotStorage'21 Presentation

July 2021

Liana V. Rodriguez, Alexis Gonzalez, Pratik Paudel, Raju Rangaswami and Jason Liu.

Learning Cache Replacement with CACHEUS | Poster Annual FIU SCIS Research Day

October 2019

Liana V. Rodriguez, Farzana Yusuf, Steven Lyons, Eysler Paz, Raju Rangaswami, Jason Liu, Ming Zhao, and Giri
Narasimhan.

ANX: Caching with Anxiety | Poster Annual FIU SCIS Research Day

October 2019

Steven Lyons, Liana V. Rodriguez, and Raju Rangaswami.

Driving Cache Replacement with ML-Based LeCaR | HotStorage'18 Poster session

July 2018

Giuseppe Vietri, Liana V. Rodriguez, Wendy A. Martinez, Steven Lyons, Jason Liu, Raju Rangaswami, Ming Zhao, and
Giri Narasimhan.

HONORS & AWARDS

USENIX Student Travel Award, FAST'18 & FAST'19 & FAST'23.

September 2018, 2019, and 2023

GAANN Fellowship from the U.S. Department of Education.

January 2022 & January 2023

CMD-IT/ACM Richard Tapia Celebration of Diversity in Computing Conference Scholarship. September 2022

Grace Hooper Celebration of Women in Computing FIU Scholarship. October 2019 & September 2022

EXPERIENCE

Team Leader | Systems Research Laboratory (SyLab), FL, US

August 2018 - December 2023

- I worked on effectively communicating the teams' goals, developing and implementing plans and strategies, and delegating responsibilities to team members. I also worked when roadblocks inevitably emerged while resolving a project problem, analyzing problems objectively, brainstorming solutions, evaluating alternatives, and implementing the optimal approach.

- I regularly monitored the team's performance, measuring their output, reviewing their progress against set targets, and providing feedback where necessary. I identified the inefficiencies and areas that required improvement by coaching team members and evaluating how the team completed tasks and achieved their goals.
- As a team leader, I helped motivate team members and worked to keep them engaged whenever I noticed a decline in their enthusiasm so that the team's productivity grew.

Graduate Research Assistant / GAANN Fellow | SyLab, FL, US

August 2017 - December 2023

- I analyzed production storage workloads and discovered workload patterns called workload primitives. As a result, I studied how these workloads function in the real world using a classification based on the composition of primitive workload types. Additionally, I computed metrics to characterize the production storage workloads. For a one-week trace from MSR, Cambridge, 3.29% of the requests were write operations, and the unique and reuse footprints were 89.24% and 2.22%, respectively. As a result, my research dives into the intricacies of the data and examines the effects of different factors, providing valuable insights. This approach has improved cache performance in novel cache replacement algorithms capable of adapting to the data and tailored solutions to meet workload needs.
- Furthermore, we used ML techniques to develop two novel, machine-learned cache replacement algorithms, LeCaR and CACHEUS, that enhance storage caching performance. When evaluating a sample trace from MSR, Cambridge, LeCaR, and CACHEUS algorithms outperform state-of-the-art caches with hit-rate values of 93.61% and 94.38%, respectively. LRU and ARC hit rates are 92.38% and 93.17%, respectively.
- Developed CaaS, a distributed caching service called CaaS that caches read and write operations and outperforms an SSD-based local cache commonly deployed in the cloud. The hit rate of CaaS is 96% for a one-day trace from Cloud VPS, outperforming the local cache hit rate of 4%. In these systems, ensuring data consistency, scalability, high availability, failover, disaster recovery, and fault tolerance must be addressed to guarantee optimal performance and reliability.

Intern | Seagate Technology, Remote, US

August 2021 - December 2021

- Seagate, a leading data storage company, has developed CORTX, an open-source object storage solution. The solution facilitates storing and managing massive amounts of unstructured data, such as media files, logs, and backups. CORTX has an integrated software stack comprising several components that work harmoniously to provide a complete object storage solution. I worked on conducting testing to ensure seamless integration and deployment of the software stack. This includes deploying cortx-motr, cortx-hare, cortx-ha, and cortx-s3server.
- I worked on CORTX's Storage Extensible Interface called File Data Manipulation Interface (FDMI), which uses plugins to implement storage features. I developed examples of plugins, such as a word count program that counts the number of times a word appears in each S3 PUT request to the storage. This plugin is equivalent to a program that polls a directory every 3 seconds and prints the 30 most popular words within each file. This research contributes to developing more robust and cutting-edge storage that can meet the ever-evolving needs of the modern world.
- I conducted a performance study on Motr, a distributed object storage system with advanced features and capabilities deployed at the Jülich Supercomputing Center (JSC) in Germany. The study aimed to analyze the system's performance and evaluate its efficiency. We reported the maximum achievable network and device bandwidth in a Sage platform x86 Motr cluster with 16 clients and 7 server nodes; each server had three tiers (tier-1, tier-2, and tier-3) that have NVMe, SSD, and HDD storage devices, respectively. The results showed 99% and 86% bandwidth utilization for read and write operations in tier-2 compared to the theoretical system limit.

Research Intern | Microsoft Research (MSR), Cambridge, UK

January 2020 - March 2020

- I worked as a member of Microsoft's Project Silica team. Specifically, my role involved helping to develop a software stack for a storage system that uses cutting-edge quartz glass technology. This innovative technology is specifically designed to cater to the requirements of Azure Cloud archival workloads, enabling efficient and cost-effective storage. Unlike traditional storage systems, Project Silica's approach involves creating tiny 'voxels' into quartz glass using a femtosecond laser, which can store vast amounts of data.
- I developed ML models to improve the storage system's software pipeline's error analysis capabilities in Project Silica, including isolation forests and encoders as anomaly detection algorithms. These algorithms can help identify unusual patterns in the data, particularly fine-grained errors within one sector (XY plane) and coarse-grained errors within one track (Z direction).
- I developed ML models that can identify anomalies when tested using real data from the Warner Bros. movie "Superman" written on glass.

Network Engineer | Telecommunications Company (ETECSA)

September 2015 - June 2016

- I maintained uninterrupted and efficient communication services, crucial for monitoring the core infrastructure of PSTN, ATM, DSL, and public and private VoIP telephone systems. This involves closely observing the hardware, software, and network elements that comprise these systems' backbone. Any potential issues were detected and resolved promptly by continuously monitoring performance metrics such as speed, latency, uptime, and number of lines interrupted.

- It was also essential to monitor the security of these systems to prevent unauthorized access and data breaches. With a focus on proactive monitoring and maintenance, these critical communication systems were highly reliable and available, ensuring seamless communication for businesses and individuals.
- I conducted a network assessment to ensure the continuous deployment of the fiber-to-the-home (FTTH) networks in Old Havana. This assessment allowed the company to identify the best action for deploying the Optical Transport Network (OTN) devices. By utilizing Gigabit Passive Optical Networking (GPON) and Ethernet Passive Optical Networking (EPON) variations of Passive Optical Networking (PON) technology, I was able to help establish and enhance network infrastructure that meets the community's needs.

Intern | Radiocommunications Company (RadioCuba)

March 2012 - July 2014

- I designed the board layout to modify the Ultra High Frequency (UHF) and Very High Frequency (VHF) communication modules from the Phase Alternate Line (PAL) to the National Television Standard Committee (NTSC) television standards. This modification improved the frequency of the television signal, resulting in a much more precise and higher-quality picture.
- I designed radiofrequency (RF) matching networks that use passive components to achieve maximum power transfer to a 50-ohm or a real-world load, such as a television (TV) antenna. This involves analyzing the properties of the RF signal and determining the best arrangement of passive components to match the load's impedance.
- I used MATLAB simulations, which modeled the results, and used virtual instrumentation of E-field sensors and spectrum/network analyzers. The simulations allowed us to visualize the RF signal's properties, including frequency, amplitude, and phase.

LEADERSHIP & ACTIVITIES

- Selected to plan and organize an event for Seagate's global CORTX Hackathon, focusing on integration challenges, possible solutions, and open-source software.
- Selected as a participant and anonymous reviewer nationwide for a peer-reviewed journal article for the ACM Transactions on Storage (TOS), the premier journal for publishing advancements in data storage research.
- Selected as a presenter for an annual Women in Computer Science (WiCS) meeting, focusing on sharing experiences and strategies for achieving academic success in computer science.
- Participated in Upsilon Pi Epsilon (UPE) activities, including workshops, presentations, tutorials, and leadership events to develop leadership skills, build industry connections, and achieve academic success.
- Selected as a volunteer for an 8-month teaching program for the Google CS First program, focusing on coding skills and learning computer science concepts using the programming language called Scratch for kids under 8.
- Selected as the team leader for four research projects at SyLab at Florida International University.

TEACHING ASSISTANT

Operating Systems

Florida International University

August 2018 & January 2019

Miami, FL

Transmission Systems II

Technological University of Havana "José Antonio Echeverría", CUJAE

2011 - 2012

Havana, Cuba

Analog Electronics I

Technological University of Havana "José Antonio Echeverría", CUJAE

2010 - 2011

Havana, Cuba

TECHNICAL SKILLS

Languages: Python, C#, C, C++, Go, R, R++, TeX, Visual Basic, High-Level Assembly (HLA).

Tools/Frameworks: GitLab, Docker, VirtualBox, VMware Workstation, VsCode, Tableau, MATLAB, LabVIEW, LaTeX, Pytorch, Scikit-learn.

CERTIFICATIONS

CompTIA A+ 220-801 Certification Exam.

CompTIA A+ 220-802 Certification Exam.

LANGUAGE

Spanish: Native.

English: Fluent and Full Professional Proficiency.

HOBBIES

Music: Traditional Cuban Music, Country Music, Pop Music, Blues, and Jazz.

Poetry: Cuban Poetry, Latino-American Poetry, Spanish Poetry, American Poetry.

Chess: Tournaments and Matches, and Titles and Rankings.

Fitness: Physical Fitness, Exercise, Training, Walking.

Animation: Computer Animation, Disney Animation, Japanese Animation.

LEADERSHIP SKILLS

Relationship Building, Emotional Intelligence, Forward-thinking, Problem-solving, Empowerment, Adaptability, Motivation, Innovation, Inspiration, Active Listening, Empathy, Presentation, Authenticity, Honesty, Creativity, Integrity, Decision-making, Accountability, Strategic Thinking, Communication, Coaching, and Management.